
The Retroactive Architecture of Agency

*On the Ontological Priority of Affective Feedback Loops
and the Phenomenology of Spurious Deliberation*

A PHILOSOPHICAL MONOGRAPH

Expanded, Annotated, and Supplemented Edition

September 2026

ABSTRACT

This monograph develops and defends a thesis concerning the causal architecture of human agency: that the primary motor of conduct is not deliberative reasoning but self-sustaining affective feedback loops, and that rationality — the apparently seamless fabric of reasons, intentions, and deliberations through which agents understand themselves — is a retroactive construction supervening upon dynamical processes it did not originate. The argument proceeds in three movements. The first diagnoses the *propositional bias* of the Western epistemic tradition: the assumption that action, to be intelligible, must be reconstructible as the output of a deliberative function, and the systematic misrecognition this produces when applied to behavior governed by attractor dynamics rather than inference. The second movement marshals the empirical warrant: the introspective-access literature initiated by Nisbett and Wilson, the split-brain experiments of Gazzaniga and colleagues and the "interpreter" hypothesis, the choice-blindness paradigm of Johansson and Hall, Wegner's theory of apparent mental causation, Haidt's social intuitionist model of moral judgment, and Mercier and Sperber's argumentative theory of reasoning. Convergent evidence from these independent programs indicates that verbal self-reports of mental causation are routinely confabulatory and that confabulation is not a pathological exception but the normal operation of a system designed to produce the appearance of reasoning. The third movement draws the ontological and intersubjective consequences: reason-giving is revealed as an interface layer rather than an engine, moral and legal accountability as a normative economy built upon ontological camouflage, and interpersonal "debate" as the collision of coupled dynamical systems masquerading as the exchange of arguments. Objections from the dual-process tradition (Paxton and Greene, Bloom) and from the compatibilist defense of conscious will (Nahmias) are considered and answered. The monograph concludes by sketching a non-propositional methodology for the human sciences: a systems-theoretic phenomenology that maps basins, gains, and thresholds where the reconstructive paradigm sought premises and conclusions.

KEYWORDS: agency; confabulation; affective dynamics; social intuitionism; apparent mental causation; attractor states; rationalization; phenomenology; dynamical systems; philosophy of action

CONTENTS

Chapter One — The Propositional Bias and the Crisis of Reconstructive Understanding	4
Chapter Two — The Dynamical Logic of Affective Feedback Loops	7
Chapter Three — The Empirical Warrant: Confabulation from the Laboratory	10
Chapter Four — Confabulation as Structural Necessity: The Retroactive Construction of Rationality	14
Chapter Five — The Phenomenology of Spurious Agency	17
Chapter Six — Intersubjective Consequences: The Rhetoric of Affect and the Illusion of the Isolated Node	19
Chapter Seven — Objections and Replies	21
Chapter Eight — Toward a Non-Propositional Theory of Human Conduct	24
Chapter Nine — Conclusion: The Seam We Cannot See	26
References	27

CHAPTER ONE

The Propositional Bias and the Crisis of Reconstructive Understanding

The Western epistemic tradition has long operated under what might be termed the *propositional bias* — the tacit assumption that human action, to be intelligible, must be reconstructible as the output of a syllogistic or decision-theoretic function. From the Aristotelian practical syllogism to the neoclassical architecture of *homo economicus*, the paradigm persists: an agent confronts a set of beliefs, weighs them against a hierarchy of desires or utilities, and emits behavior as the logical consequent of this internal computation. Understanding, on this model, is fundamentally *reconstructive*; to comprehend an action is to simulate the reasoning that produced it, to inhabit, however vicariously, the inferential space of the actor.

The genealogy of this assumption is worth tracing, for its depth explains its invisibility. In the *Nicomachean Ethics*, Aristotle models practical action on the structure of deduction: a major premise articulating a universal end (the good to be achieved), a minor premise registering a particular fact (this action, here, achieves it), and a conclusion that is not a proposition but a deed.¹ The schema is seductive because it converts the opacity of motivation into the transparency of proof. To know why the agent acted is to possess the syllogism from which the act follows. Two millennia later, the same schema reappears in secularized form: Davidson's doctrine that a reason rationalizes an action by being its cause gives philosophical dignity to the folk practice of reason-attribution,^[1] while the expected-utility framework of the economists transforms the syllogism into a maximization problem without altering its essential grammar — premises in, conduct out.

What unites these otherwise heterogeneous traditions is a structural commitment, not a doctrine. The commitment is this: *the causal order and the justificatory order coincide*. The reasons that would justify the action are identical with (or at least isomorphic to) the causes that produced it. Given this identity, understanding another agent is an exercise in reverse engineering: observe the output, reconstruct the computation, and the behavior becomes legible. The reconstructive paradigm is thus not merely a theory of action but a theory of *intersubjectivity* — it specifies what it is to make sense of one another.

1. Aristotle's practical syllogism is deliberately asymmetrical with theoretical syllogism: where the latter terminates in assent, the former terminates in motion. Yet the formalism is the same — the deed is the conclusion, and hence, in principle, derivable. The deep claim is that action without premises is not action at all but mere movement.

1.1 The Limits of Reconstruction

Yet the paradigm encounters a systematic failure when confronted with a class of behaviors that resist such reconstructive assimilation — not because they are random or chaotic, but because they are governed by a logic that is not propositional but dynamical. The violent offender in the moment of escalation; the subject seized by panic, by infatuation, by moral revulsion; the crowd swept into collective frenzy: these figures do not fail to be understood because they are irrational in the sense of reasoning badly. Their conduct does not exhibit the failure modes of inference — invalid transitions, contradictions between premises, misweighted evidence. It exhibits, instead, the signature of a different kind of system altogether: positive feedback, threshold crossing, basin capture. The grammar of their conduct is not that of reasoning at all. It is the grammar of loops.

The failure of reconstruction here is principled, not practical. Reconstruction presupposes that there exists, somewhere in the causal ancestry of the act, a set of propositionally structured states playing the role of premises. Where no such states exist — where the causal work is done by a self-amplifying dynamical process whose states are not propositional and whose transitions are not inferential — the reconstructive method does not merely yield a poor approximation. It yields a confident falsehood: a fabricated syllogism dressed in the grammar of reasons, attributed to a system that never entertained it. The error, then, lies not in the subject but in the analyst's insistence that all behavior must be legible as the conclusion of an argument.

1.2 From Reconstruction Failure to Systematic Confabulation

What emerges from this misalignment is a more profound and structurally pervasive phenomenon, and it is the central phenomenon of this monograph: the *retroactive imposition of rationality* upon conduct that was never rational to begin with. Human beings do not merely act on emotion and later acknowledge this fact; they act on emotion and subsequently confabulate a rational architecture that renders the action intelligible within the normative economy of reasons-giving. The analyst is not the only one who fabricates syllogisms. The subject fabricates them too — fluently, sincerely, and with a conviction that admits of no first-person doubt.

This is not mere dishonesty, nor is it a failure of introspective access in the narrow sense — as though the true reasons were stored in some inner archive to which the subject had unfortunately misplaced the key. It is, rather, a constitutive feature of self-conscious agency: the mechanism by which a creature governed by affective feedback loops maintains the phenomenological illusion of being a creature governed by reasons. The claim to be defended is therefore stronger than the familiar psychological observation that

self-reports are unreliable. It is an ontological thesis about the *order of operations* in the generation of agency: the loop is primary; the reason is derivative; and the derivative systematically masquerades as the primary.

THESIS (RETROACTIVE ARCHITECTURE)

For a large and systematically under-acknowledged class of human actions, the propositional content retrospectively cited as the action's reason was not causally operative in its production. The reason is generated after (or in parallel with) the action by a distinct interpretive process whose function is not control but narration. The experienced priority of reasons over actions is an artifact of this architecture.

The chapters that follow develop this thesis along three axes: the dynamical character of the affective substrate (Chapter Two), the empirical evidence for the interpretive superstructure (Chapters Three through Five), and the intersubjective and methodological consequences (Chapters Six through Eight). But the polemical point of departure should be stated plainly. The reconstructive paradigm has not failed because we have not yet found the right model of deliberation. It has failed because it mistakes the direction of fit between reasons and actions — assuming that reasons flow into actions, when the dominant flow runs the other way.

CHAPTER TWO

The Dynamical Logic of Affective Feedback Loops

To grasp the ontological priority of emotional feedback loops over deliberative reasoning, one must first abandon the folk-psychological picture of emotions as transient perturbations of an otherwise rational cognitive system. On that picture — which survives, barely disguised, in much of the cognitive-scientific literature as the view that emotion "biases" or "modulates" cognition — the agent is a reasoning engine occasionally knocked off course by storms of affect, storms that pass and leave the engine intact. This picture is doubly wrong. Emotions are not interruptions; they are control structures. And they are not passengers on a rational vehicle; on the view defended here, they *are* the vehicle, and reason is the passenger who writes the travelogue.

Emotions operate not as inputs to a central decision processor but as self-sustaining dynamical systems characterized by positive feedback, variable gain, and attractor states. The dynamical-systems approach to emotion — developed most rigorously in the work of Lewis and Granic and their collaborators — models emotional episodes as emergent, self-organizing configurations of appraisal, physiology, attention, and action-readiness, stabilized by recursive coupling among their components rather than orchestrated by any central controller.^{[2][3]} What the phenomenologist calls "being in the grip of an emotion" the systems theorist calls capture by an attractor: a region of the system's state space toward which trajectories converge and from which they escape only under perturbation exceeding some threshold.

2.1 The Anatomy of a Loop

Consider the architecture of rage. An initial perturbation — a perceived slight, a territorial intrusion, a status injury — triggers a cascade of physiological and neurochemical events: catecholamine release, cardiovascular recruitment, amygdala-driven reprioritization of perceptual processing. These events do two things at once. They lower the threshold for aggressive response, and they narrow the field of perceptual attention to cues that confirm the appropriateness of that response. The narrowing is crucial, for it is the recursive term in the loop: the angry subject does not merely feel rage and see the world; the rage selects what counts as the world — the remembered grievances that surface unbidden, the ambiguous gesture now read as contemptuous, the voice that sounds, in retrospect, to have been asking for it. The selected world then feeds back as fresh confirmation of the rage, which intensifies, which selects further, and so on around the cycle.

The subject does not decide to be angry; the anger decides for him — and in deciding, it restructures the informational environment such that the anger appears, from the inside, as a proportionate and inevitable reaction to objective features of the world. This last

clause deserves emphasis. The phenomenology of the loop is not a phenomenology of compulsion. The subject in an attractor does not feel the walls of the basin; he feels the solidity of the ground. The anger presents itself with the same evidentiary dress as perception: *this is what happened; this is what he did to me; this is what he deserves.*

2.2 Loop and Argument: A Formal Contrast

The contrast between the two architectures can be stated with some precision. In a logical system, premises stand in relations of justification to conclusions; the flow is acyclic, directed, and externally evaluable. One can audit an argument: expose its premises to evidence, check its transitions for validity, suspend its conclusion. The auditing stance is possible precisely because the argument is not self-feeding — changing the premises changes the conclusion, and the premises themselves are not produced by the conclusion they support.

In a feedback loop, by contrast, the output is recursively fed back into the input, amplifying the system's deviation from equilibrium until a threshold is crossed or an external constraint interrupts the cycle. Schematically, for an affective state variable $A(t)$ coupled to a perception–appraisal process P : the appraisal generates affect, and the affect re-parameterizes the appraisal, so that P becomes a function of A even as A remains a function of P . The system has no premises in the logical sense; it has phases and basins. It does not justify; it propagates. And its "evaluation" cannot be conducted from the inside, because the evaluating faculty — the appraisal process that would have to stand in judgment — is itself one of the coupled variables.

Table 1. Structural contrast between inferential and dynamical control of conduct

DIMENSION	INFERENTIAL ARCHITECTURE (ARGUMENT)	DYNAMICAL ARCHITECTURE (LOOP)
Basic units	Propositions (beliefs, desires, utilities)	Coupled state variables (arousal, attention, appraisal, motor readiness)
Transition rule	Inference; revision by evidence	Iteration; amplification by feedback
Flow structure	Acyclic, premise → conclusion	Recurrent, output → input
Stability	Consistency under scrutiny	Attractor capture; convergence to basin
Mode of change	Counter-evidence, better arguments	Perturbation, threshold crossing, external constraint
First-person appearance	"I concluded that..."	"I saw how things were..."

Failure mode	Invalidity, inconsistency	Runaway escalation, hysteresis, entrainment
Auditability	Externally evaluable in principle	Self-sealing; evaluation faculty inside the coupling

2.3 The World-Structuring Effect

Because the loop is self-sustaining, it generates what phenomenologists might call a *world-structuring effect*: the subject does not experience himself as caught in a loop but as responding to how things are. The feedback is invisible from the first-person perspective because the loop itself constitutes the perspective. This is the deepest point of contrast with the older vocabulary of "passion." The tradition from Plato's charioteer to Hume's calm and violent passions imagined emotion as a force acting *upon* a self that stands apart from it — a self that can, at least in principle, regain the reins. The dynamical picture is less comforting. There is no charioteer waiting to resume control, because the perspective from which control would be exercised is itself a configuration of the system. What returns, when the episode subsides, is not a suppressed sovereign but a different attractor — the "calm" state is another basin, not the absence of basins.

The implications for agency are severe, and they set the agenda for everything that follows. If the primary motor of human conduct is not deliberation but dynamical entrainment, then the traditional vocabulary of choice, decision, and intention applies not to the generative phase of action but to its post-hoc narrativization. The action emerges from the loop; the reason emerges from the action. This inversion of the standard causal order is not incidental but systematic — and its systematicity is exactly what the next three chapters document.

CHAPTER THREE

The Empirical Warrant: Confabulation from the Laboratory

Philosophical theses about the architecture of agency live or die by their empirical plausibility. The present chapter therefore pauses the argument to assemble the evidentiary base — five research programs, largely independent of one another, converging on a single structural finding: human beings produce confident, fluent, and false accounts of the causes of their own judgments and actions, and they do so not occasionally but as a matter of routine.

3.1 The Introspective-Access Result

The foundational document is Nisbett and Wilson's 1977 review, *Telling More Than We Can Know*.^[4] Surveying dozens of studies in social psychology, they demonstrated that subjects' verbal reports about the stimuli that influenced their judgments were frequently and demonstrably wrong — wrong in both directions, attributing influence to inert factors and denying influence to controlling ones. In their most famous demonstration, passersby asked to evaluate identical pairs of stockings arrayed on a table showed a strong position effect, overwhelmingly preferring the rightmost item; yet no subject cited position as a factor, and when it was suggested, subjects dismissed it as implausible, offering instead confident references to texture, color, and workmanship. Nisbett and Wilson's conclusion was not modest: there may be, they wrote, "little or no direct introspective access to higher order cognitive processes." When people report on the causes of their judgments, they do not read off an internal record; they consult culturally supplied theories of what plausibly causes judgments of that kind. The report is an inference, not an observation.

The philosophical weight of this result is easy to understate, because it is so often framed as a finding about the fallibility of memory or the haziness of self-perception. It is neither. The reports in question are produced *immediately*, with the judgment still warm; their errors are not decay but construction. What Nisbett and Wilson uncovered is that the functional architecture of self-explanation is generative: the mind manufactures accounts where an audit trail would be, and the manufactured account is subjectively indistinguishable from recollection.

3.2 The Interpreter: Confabulation at the Neural Level

Gazzaniga's split-brain experiments supplied the phenomenon with a mechanism and a name.^{[5][6]} In patients whose corpus callosum had been severed, stimuli presented to the right hemisphere (via the left visual field) could drive behavior without any verbal

awareness: shown a snow scene in the left field and a chicken claw in the right, a patient asked to select associated pictures pointed with the left hand to a shovel and with the right to a chicken — and then, asked to explain, the speaking left hemisphere responded without hesitation: "The chicken claw goes with the chicken, and you need a shovel to clean out the chicken shed." The explanation is fluent, internally coherent, and entirely fabricated. The left hemisphere observed the behavior and narrated it, integrating the unexplained action into what Gazzaniga calls the ongoing storyline through which the brain makes "all sensible" what the organism is doing.^[7]

Two features of the interpreter deserve emphasis. First, its operation is *instantaneous and unhesitating* — there is no phenomenology of confabulation, no felt gap between the action and its story. Second, the interpreter is not an artifact of commissurotomy. The surgical cases make the mechanism visible by severing the information channel, but the mechanism itself — an explanatory module that reconciles ongoing behavior with the self's running narrative — is standard equipment. In intact brains, the channel is open and the interpreter's fabrications are correspondingly harder to catch; the split-brain patient is the control condition for a device we all carry. Gazzaniga's own conclusion is unguarded: perhaps "at the heart of consciousness, it's a story that we've told ourselves. And it's ongoing."^[7]

3.3 Choice Blindness: Confabulation in Real Time

The choice-blindness paradigm of Johansson, Hall, and colleagues brought the phenomenon under experimental control in normal subjects.^[8] Participants shown pairs of faces and asked which they found more attractive were, on manipulated trials, handed the photograph they had *rejected* and asked to explain their choice. A large majority failed to detect the switch — and proceeded to offer detailed, introspectively derived reasons for a choice they never made: "She had earrings," said one subject of a face without them. Crucially, subsequent analysis of the verbal reports found almost no systematic difference between manipulated and non-manipulated trials: the confabulated reasons were linguistically and semantically indistinguishable from genuine ones.^[9] The implication cuts deep. If fabricated justifications and veridical ones are produced by the same process with the same surface properties, then the sincerity and fluency of a reason-report carry no evidential weight whatsoever about its causal accuracy.

3.4 Apparent Mental Causation and the Illusion of Will

Wegner generalized the finding into a theory of the experience of agency itself.^{[10][11]} On his theory of *apparent mental causation*, the feeling of conscious will is not a readout of the causal sequence producing an action but an inference — drawn whenever a thought about the action precedes it, is consistent with it, and is unaccompanied by salient rival causes.

Because the inference can be fooled — as in Ouija-board spelling, facilitated communication, hypnosis, and spirit possession, where actions occur without experienced will, or experienced will without action — the experience of authorship is decoupled from the fact of authorship. "Experiences of conscious will," Wegner concludes, "arise from processes whereby the mind interprets itself — not from processes whereby mind creates action."^[11] The subject's felt priority — *I thought, then I acted* — is the temporal artifact of an interpretive process that arrives on the scene after the causal work is underway and backdates its own contribution.

3.5 Moral Dumbfounding and the Rational Tail

Finally, Haidt's social intuitionist model extended the confabulation thesis into the moral domain — the domain where the stakes of reason-giving are highest.^[12] Reviewing evidence from social, cultural, and biological psychology, Haidt argued that moral judgment is generally the product of fast, automatic intuitions, with moral reasoning following as "a post hoc construction, generated after a judgment has been reached." His signature experimental phenomenon is *moral dumbfounding*: subjects confronted with harmless-but-taboo scenarios maintain their condemnation with full confidence even as every argument they offer is dismantled — and rather than revising the judgment, they search for new arguments or simply assert "I just know it's wrong." The judgment is the fixed point; the reasons are the variable. Haidt's formulation — moral reasoning is "the emotional dog's rational tail" — captures precisely the causal order this monograph defends: the affective system settles first, and the rational apparatus is subsequently mobilized as its press secretary.

3.6 Convergence

Table 2. Five convergent research programs on post-hoc self-explanation

PROGRAM	PARADIGM	CORE FINDING	IMPLICATION
Nisbett & Wilson (1977)	Verbal reports on judgment	Reports causally inaccurate; based on implicit theories, not access	Self-explanation is inferential, not observational
Gazzaniga et al. (1970s–2000s)	Split-brain experiments	Left hemisphere instantly narrates behavior it did not cause	Dedicated interpretive module; fabrication is fluent and unfelt
Johansson & Hall (2005–)	Choice blindness	Subjects justify choices they never made; reports from indistinguishable genuine ones	Sincerity and fluency are no evidence of causal accuracy

Wegner (1999–2002)	Anomalies of will; automatisms	Experience of will is an inference from priority, consistency, exclusivity	Felt authorship decoupled from actual authorship
Haidt (2001)	Moral dumbfounding	Moral judgments survive the refutation of every supporting argument	Reasoning is post-hoc justification of intuitive verdicts

The convergence is the argument. Five programs, different methods, different populations — surgical patients, passersby, laboratory subjects, moral judges — one structure: a dynamical process produces the output, and an interpretive process produces the account, and the two are related not as cause and readout but as event and press release. The remaining question is not whether confabulation occurs but *why it exists at all* — why a creature would evolve, or be acculturated into, an architecture in which the narrator claims the driver's seat. That is the question of Chapter Four.

CHAPTER FOUR

Confabulation as Structural Necessity: The Retroactive Construction of Rationality

The cognitive science literature on confabulation has long documented the human propensity to generate explanatory narratives for behaviors whose actual causes are opaque to introspection. But the standard framing of this literature — as a catalogue of illusions, errors, and failures of self-knowledge — misses the deeper point, and missing it, misses the architecture. Confabulation is not a failure of a system designed to report its own reasoning; it is the normal operation of a system designed to produce the *appearance* of reasoning. Once this is seen, the question shifts from "why do our self-reports err?" to "what is the functional role of a faculty whose characteristic product is a false-but-convincing causal story?"

4.1 The Layered Architecture of the Self

The human subject is not, first and foremost, a reasoner who occasionally lapses into affect. The human subject is an affective system that has evolved — or been acculturated — to produce reasons as a byproduct of action. This is not to reduce all human conduct to mere reactivity; it is to insist that the layer of rationality is a *supervening* layer, one that arrives after the fact and claims authorship of events it did not originate. The loop fires; the body moves; and then the left-hemisphere interpreter, that tireless fabulist, weaves a narrative in which the movement was the conclusion of a practical syllogism.²

The subject experiences this narrative not as a story told about himself but as the very structure of his agency. He was angry, yes, but he was angry *because of* the insult, and the retaliation was proportionate, and the alternative courses of action were considered and rejected. The loop becomes a line of reasoning. The basin of attraction becomes a weighing of reasons. Note the specificity of the transformation: it is not that the subject vaguely misremembers what happened. It is that the output of the interpretive process has the *form* of practical reasoning — premise, deliberation, proportionate response, considered alternatives — because that is the form the normative environment demands. The confabulation is shaped by its audience, even when the audience is only the self.

2. Hirstein's neuropsychological synthesis reaches the same architectural conclusion from clinical data: confabulation, across its many syndromes, is not disordered storytelling but the unmasking of the brain's ordinary knowledge-fixation and narration routines operating with damaged error-checking. The pathology reveals the physiology. See Hirstein (2005), *Brain Fiction*.^[13]

4.2 Why the Faculty Exists: Reasoning as Argumentation

What function could such a faculty serve? The most developed answer in the contemporary literature is Mercier and Sperber's argumentative theory of reasoning.^[14] On their account, the evolved function of human reasoning is not the solitary improvement of knowledge and decision but the production and evaluation of arguments in social exchange — reasoning as a device for persuasion and for epistemic vigilance toward the persuasion of others. The theory explains, rather than laments, the signature "defects" of human inference: the confirmation bias, the myside bias, the tendency of solitary reasoning toward confident error. These are malfunctions only if reasoning's function is truth-tracking within a single skull; they are design features if its function is to win assent and to defend positions already held.

The argumentative theory supplies the missing link in the present argument. The confabulatory faculty is not a broken reporter; it is a working press office. Its product — a fluent, premise-ordered, audience-calibrated justification — is exactly what a social organism needs when the currency of coordination is reasons. The subject who can instantly produce a syllogism for any action can negotiate, can persuade, can be held to account and hold others to account, can accumulate reputation. The subject who could only say "my affective system crossed a threshold" would be socially illegible — would fail, precisely, to qualify as an agent in the only sense the collective recognizes.

4.3 The Social Function of Ontological Camouflage

This is the deeper sense in which retroactive rationalization is structurally necessary rather than merely habitual. Human intersubjectivity is organized around the exchange of reasons. We hold one another accountable not for our loops but for our justifications. The moral and legal order presupposes that actions are the termini of deliberative processes; to be a person — a rights-bearer, a promise-maker, a defendant — is to be capable of offering and accepting reasons. The confabulatory mechanism thus operates as a kind of *ontological camouflage*, allowing affect-driven creatures to participate in a normative economy that requires the simulation of rational agency. We do not merely rationalize to save face; we rationalize to maintain the structural integrity of the social world.

The camouflage is ontological, not merely strategic, because it runs all the way down. There is no backstage self who knows the truth and feeds lines to the performer. The interpreter's product is consumed first and most thoroughly by its producer: the subject believes his own press releases, and must, for the performance to be stable under interrogation. Wegner notes the moral dimension of this arrangement with characteristic precision: the experience of conscious will, whatever its causal impotence, functions as "the person's guide to his or her own moral responsibility for action" — the internal marker that tells the

organism which events to claim and which to disown.^[15] The narrative is load-bearing. A society of transparent dynamical systems would need a completely different theory of blame.

4.4 Confabulation Is Not Lying

A clarification is necessary, because the vocabulary of fabrication invites a misreading. The liar knows the truth and conceals it; the confabulator has no truth to conceal. The actual causal process — the loop, the basin, the threshold — is not propositionally encoded anywhere in the system; there is nothing to "access," correctly or incorrectly. This is why the phenomenology of confabulation is the phenomenology of candor, and why introspective training, honesty, and intelligence provide so little protection. The defense of the position must therefore be structural, not moral: one cannot ask the subject to stop confabulating, only to notice the seam where the story meets the loop. What such noticing might look like — and whether it is even coherent as a project — is taken up in Chapter Eight.

CHAPTER FIVE

The Phenomenology of Spurious Agency

From the first-person perspective, the retroactive imposition of rationality is not experienced as imposition at all. The subject does not think, *I acted on impulse, but I will now invent a reason*. He thinks, *I acted for this reason* — and the temporal order of the experience is such that the reason appears to precede and to cause the action. This is the phenomenology of spurious agency: the lived experience of being the author of one's conduct when one is, in fact, the epiphenomenon of a dynamical process.

5.1 The Backdating of Authorship

The depth of this illusion is difficult to overstate, and its mechanism is worth specifying. Wegner's three criteria for the experience of will — priority, consistency, exclusivity — describe an inferential engine that assigns authorship whenever a suitable thought is found in temporal proximity to an action.^[10] But the engine's output is not a proposition (*my thought probably caused that*); it is a *feeling*, the feeling of doing, and feelings do not present themselves as inferences. The result is a systematic temporal illusion: the interpretive process, which necessarily follows the initiation of the causal cascade, experiences its own product as antecedent. The press release is dated earlier than the event it reports. Because the narration is seamless and immediate, there is no experienced gap into which doubt could insert itself — the subject never catches the interpreter in the act, because catching it would require a perspective outside the very narration that constitutes his perspective.

Consider, then, the experience of making up one's mind. The phenomenology is one of internal deliberation: options are entertained, consequences projected, values weighed. The phenomenologist's description is accurate as description — these events do occur in the stream of consciousness. Yet the empirical record suggests that in many cases, the outcome of this deliberation is determined by affective valence before the deliberation begins. Damasio's somatic marker hypothesis provides the physiological sketch: biasing signals, generated covertly from body-state representations, constrain the option space and weight the alternatives before any explicit reasoning gets underway — the "reasoning" operates on a shortlist that feeling has already drawn up.^[16] The subject experiences himself as arriving at a conclusion through reasoning; what has actually occurred is that the feedback loop has settled into an attractor state, and the deliberative process has served as a post-hoc validation mechanism — a kind of cognitive security clearance granted to a decision already made by the body. The reasoning is real, in the sense that it occurs; but it is not causally efficacious in the way the subject assumes. It is the shadow of the loop, not its source.

5.2 The Recursive Deepening of the Basin

This structure is recursive, and the recursion is what converts episodes into character. The confabulated reason, once articulated, becomes available as a premise in future deliberations. The subject who has justified his anger as a response to injustice does not merely possess a cover story; he possesses a standing reason, a node in his self-concept that will prime similar emotional responses in the future — lowering the threshold for the next capture, pre-fabricating the next justification. The loop generates the reason, and the reason feeds back into the loop, deepening the basin of attraction. Over time, the subject becomes a character in his own narrative: a creature of principle and conviction whose emotional outbursts are always, somehow, principled; whose appetites are always, somehow, justified preferences; whose aversions are always, somehow, considered positions. The mask becomes the face. The retroactive rationalization becomes the proactive identity.

There is a further, subtler consequence. Because each new justification is woven from the materials of previous ones, the narrative layer acquires a kind of internal integrity — consistency, coherence, even a style — that is entirely genuine as a property of the *story* while bearing only a contingent relation to the dynamics of the loops it chronicles. The subject's self-model can be a masterpiece of novelistic construction and a garbage model of his own causation at the same time. This explains a phenomenon that purely psychological accounts find puzzling: the extraordinary resistance of self-narratives to corrective evidence. The narrative is not an empirical hypothesis about the loop, held tentatively and revised on evidence; it is the fabric in which evidence is interpreted. One does not revise the fabric; one re-interprets the anomaly.

5.3 Two Orders of Description

The upshot can be stated as a distinction between two orders of description, both valid, neither reducible to the other. In the first order — the order of the loop — the subject is a dynamical system: coupled variables, basins, thresholds, gains. In the second order — the order of the narrative — the subject is an agent: beliefs, desires, intentions, reasons. The spuriousness of agency consists not in the second order being false but in its systematic misrepresentation of its own causal status: the narrative claims to *drive* what it in fact *chronicles*. And because social life is conducted entirely in the second order — contracts, promises, accusations, excuses, all trade in reasons — the misrepresentation is not an optional error the subject might correct. It is the price of admission to the human conversation.

CHAPTER SIX

Intersubjective Consequences: The Rhetoric of Affect and the Illusion of the Isolated Node

The social and political consequences of the architecture described so far are profound, and they begin with a deflation of our most self-congratulatory institution: reasoned public discourse. Intersubjective discourse — moral argument, political debate, legal adjudication — presupposes that the participants are engaged in the exchange of genuine reasons, that the surface structure of argumentation maps onto a deep structure of deliberation. But if the primary motor of human conduct is the affective feedback loop, then much of what passes for rational discourse is in fact the collision of loops masquerading as reasons.

6.1 Debate as Dynamical Coupling

Two subjects, each caught in a self-sustaining emotional basin, each generating post-hoc rationalizations that render their own position as the necessary conclusion of sound reasoning, engage in what appears to be a debate. They trade arguments, invoke principles, appeal to evidence. But the causal structure of the encounter is not that of two rational agents converging on truth through the force of the better argument. It is the dynamical interaction of two feedback systems, each of which interprets the other's outputs as perturbations to be either absorbed or amplified. The arguments are not the control signal; they are the surface noise of deeper affective entrainment. The debate is not won by the side with better logic; it is won by the side whose loop has the higher gain, whose basin is deeper, whose confabulatory apparatus is more robust.

This description is not cynicism; it is supported by the very features of reasoning that the argumentative theory predicts. If reasoning evolved for persuasion rather than truth-tracking,^[14] then we should expect exactly what we observe: myside bias that strengthens under challenge, the backfire dynamics of corrective information, the phenomenon — documented across Haidt's studies — of positions held more confidently after their supporting arguments have been destroyed.^[12] A truth-tracking system retreats under refutation; a justification-generating system escalates. The observed behavior of disputants is the behavior of the second kind of system.

6.2 The Illusion of the Isolated Node

Beneath the fantasy of rational debate lies a deeper structural error: the illusion of the isolated node — the assumption that individual agents are independent reasoning units whose beliefs and desires can be evaluated in isolation from the dynamical systems in

which they are embedded. In reality, agents are nodes in a network of affective contagion, their ostensibly rational outputs generated by loops that are coupled, recursively, to the loops of others. Emotional states propagate through groups with the mechanics of coupled oscillators: synchrony, entrainment, resonance. The crowd is not a collection of reasoners who happen to agree; it is a phase-locked dynamical system whose individual members' "opinions" are local expressions of a collective attractor.

Once this is granted, a series of invidious distinctions collapses. The violent criminal and the righteous ideologue are not opposites on a spectrum of rationality; they are variations on the same underlying architecture, differing only in the content of their confabulations and the social acceptability of their threshold crossings. The respectable partisan whose loop drives him to write furious columns and the offender whose loop drives him to throw punches share a generative grammar; what separates them is the coupling of their basins to different social enforcement structures. This is not moral equivalence — the outputs differ, and the outputs matter enormously — but it is architectural identity, and mistaking the difference for one of *kind* rather than *degree and coupling* is precisely the misrecognition the propositional bias enforces. The law courts and the opinion pages are, from the dynamical standpoint, the same phenomenon wearing different costumes: arenas in which loop-generated conduct is adjudicated in the currency of reasons.

6.3 The Accountability Interface

None of this implies that the reason-exchanging layer is dispensable. On the contrary: the interface is where coordination happens, and a normative economy built on reasons is — despite its foundations in camouflage — a genuine civilizational achievement, the machinery by which affect-driven primates manage cooperation at scale. The point is diagnostic, not abolitionist. When we understand that the interface is an interface, we stop being surprised by its characteristic failures: the futility of fact-checking against basin-deep conviction, the escalation of argument into polarization, the way "rational discourse" so reliably tracks alliance rather than evidence. And we begin to ask the engineering questions that the propositional bias forbids: not "how do we get people to reason better?" but "how do we alter the coupling, the gains, the basins?" — questions about the design of environments, incentives, and rhythms rather than the improvement of arguments.

CHAPTER SEVEN

Objections and Replies

A thesis as sweeping as the retroactive architecture invites principled resistance, and it deserves the strongest available versions of it. Three objections, drawn from the philosophical and psychological literature, are considered here; a fourth, deeper worry — the threat of incoherence — closes the chapter.

7.1 The Compatibilist Defense of Conscious Will

Nahmias, reviewing Wegner, articulates the most pointed version of the objection.^[15] The evidence for confabulation, he argues, consists of *exceptions* — split-brain patients, Ouija boards, hypnotic suggestion, position effects on stocking choice — and from the existence of exceptions nothing follows about the rule. That conscious will can be dissociated from action in contrived cases does not show that, in the normal case of correspondence, it is causally inert: "just because A can occur without B does not mean that when A does occur with B, A is not a cause of B." Moreover, conscious deliberation has demonstrable downstream effects — in planning, in the formation of distal intentions, in the training of habits that later run without felt will (the athlete, the musician). The interpreter may embroider, but the self that plans is no epiphenomenon.

Reply. The objection is sound as far as it goes, and the present thesis should be stated to accommodate it. Nothing here denies that consciously entertained reasons can be causally efficacious — in planning horizons, in pre-commitment, in the slow sculpting of the basins themselves. The claim is about the *default causal order in the generative moment of action*, and about the systematic misrepresentation of that order in experience. Nahmias's examples are instructive precisely because they show where deliberation bites: upstream of the loop, in the design and training of the dynamics — exactly where a dynamical theory predicts that propositional content could matter. What the compatibilist defense does not touch is the retroactive thesis itself: that whatever role conscious reasons play, the experience of their causal priority is manufactured after the fact, and that this manufacturing is indistinguishable, from the inside, from genuine authorship. The objection saves the efficacy of planning; it does not save the phenomenology.

7.2 The Dual-Process Rejoinder

From moral psychology, Paxton and Greene press the case that reasoning plays a substantive role in moral judgment — counteracting automatic bias, overriding intuition, driving the distinctly "utilitarian" verdicts that intuition resists.^[17] Bloom adds the historical argument: intuition alone cannot explain moral change — abolition, emancipation, the expanding circle — which seems to require exactly the kind of

argument-driven reconsideration that the intuitionist picture treats as epiphenomenal.^[18] If reasoning were mere rationalization, moral revolutions would be inexplicable.

Reply. The dual-process evidence constrains the thesis but does not overturn it; indeed, the dynamical framework predicts the very phenomena adduced. Controlled reasoning, on the present account, is not an alternative to the loop but an *external constraint upon it* — a source of perturbation capable, when sufficiently strong and repeatedly applied, of driving the system out of one basin and into another. Moral revolutions, accordingly, are not counterexamples but slow-motion basin restructuring: arguments do not replace intuitions by force of logic; they propagate through populations as coupled perturbations, gradually reshaping the affective landscape until new intuitions — new fixed points — condense. The abolitionist pamphlet did not win a debate inside individual skulls; it changed what those skulls found revolting. Haidt himself grants this channel explicitly: interpersonal discussion can activate new intuitions, which then carry forward into future judgments.^[19] Reasoning matters — but as an input to the dynamical system, not as its engine, and its causal route runs through affect even when its content is logic.

7.3 The Charge of Self-Refutation

If all reason-giving is post-hoc confabulation, then this monograph's own arguments are confabulations, and its thesis self-defeating. The objection is venerable — it shadows every deflationary account of reason from Nietzsche to the evolutionary debunkers — and it deserves better than dismissal.

Reply. The thesis is not that reasoned argument is impossible or uniformly impotent, but that its causal role is misdescribed by the reconstructive paradigm. The argumentative theory, recall, predicts that reasoning is *better* in its proper element — the evaluation of others' arguments under conditions of genuine social friction — than in solitary self-justification.^[14] Scholarship, at its best, institutionalizes exactly this corrective coupling: a machinery for subjecting loop-generated convictions to adversarial perturbation. The present argument does not exempt itself; it aspires to participate in the one practice — intersubjective criticism — the architecture itself identifies as the only reliable check on confabulation. The seam cannot be seen from inside one's own loop; that is why we read each other.

7.4 The Quiet Worry: Eliminativism

Finally, there is the worry that the whole apparatus — loops, basins, interpreters — is a eliminativist reduction that dissolves the person into neurochemistry and thereby loses the very phenomena (guilt, love, conviction) that any adequate philosophy of agency must save. The reply has been built into the argument from the start: the narrative layer is real, causally significant, and indispensable. What is denied is only its self-ascribed causal

priority. To say that the story is a story is not to say that it is nothing — it is to say what kind of thing it is, and hence what can be asked of it.

CHAPTER EIGHT

Toward a Non-Propositional Theory of Human Conduct

What would it mean to take seriously the ontological priority of affective feedback loops? It would require, first, a fundamental reorientation in our approach to understanding human action. We would need to abandon the reconstructive paradigm — the assumption that to understand another is to simulate their reasoning — and adopt instead an explanatory paradigm that treats behavior as the observable output of a dynamical system whose internal states are not necessarily accessible to propositional articulation.

8.1 A Methodological Inversion

Understanding the violent offender would not mean trying to think as he thinks, for there may be no thinking, in the deliberative sense, to be reconstructed. It would mean mapping the topology of his emotional basins, the gain of his feedback loops, the thresholds at which his system crosses into action. The shift is from hermeneutics to systems analysis: from *what were his reasons?* to *what are his attractors, what perturbs him across them, what is the coupling between his appraisal processes and his arousal?* The questions sound clinical, and they are — but the claim of this monograph is that they are the general form of all understanding of human conduct, of which the reconstructive method is a special case applicable only where the loop's dynamics happen to have been pre-structured by genuinely deliberative processes (planning, pre-commitment, institutional constraint).

This is not to embrace a crude eliminativism that reduces human agency to neurochemistry or mechanics. The confabulatory layer — the layer of reasons, narratives, and justifications — is a real and causally significant feature of human life. But it is a secondary feature, one that operates not as the engine of conduct but as its *interface*. The interface is where we hold one another accountable, where we construct moral and legal orders, where we tell the stories that make collective life possible. But to mistake the interface for the engine is to commit a category error with far-reaching consequences. It leads us to debate the logic of arguments that were never the actual control signal, to hold subjects responsible for deliberations they never performed, and to invest in the rational modification of beliefs when what requires modification is the architecture of the loop itself.

8.2 Consequences for Practice

The inversion has concrete implications across the domains where the propositional bias currently governs institutional design.

Table 3. The reconstructive and dynamical paradigms as guides to practice

DOMAIN	RECONSTRUCTIVE PARADIGM (CURRENT DEFAULT)	DYNAMICAL PARADIGM (PROPOSED)
Understanding others	Simulate their reasoning; reconstruct the syllogism	Map basins, gains, thresholds; identify coupling and contagion
Criminal justice	Assess deliberation and intent at the moment of action	Assess the dynamics of escalation; intervene on thresholds and environments
Public discourse	Improve arguments; correct misinformation; fact-check	Redesign coupling and incentives; alter the rhythms and gains of propagation
Psychotherapy	Insight through verbal reconstruction of reasons	Basin restructuring: exposure, state regulation, rehearsal under perturbation
Self-knowledge	Introspection; "know thy reasons"	Suspicion of fluent self-report; tracking of one's own triggers, states, and escalations
Ethics education	Teaching principles and moral reasoning	Cultivating intuitions and habits; designing situations that train the loop

8.3 The Task

The task, then, is not to eliminate the loop in favor of reason, nor to celebrate the loop as a more authentic mode of being. Both moves are nostalgic — the first for a rational sovereign that never existed, the second for a pre-reflective innocence that language has forever foreclosed. The task is to recognize the bifurcated nature of human agency: a creature of feedback who lives as a creature of reasons, a dynamical system who experiences himself as the author of arguments he never made. To understand such a creature requires not the sympathy of shared reasoning but the rigor of systems thinking. It requires the humility to acknowledge that some behaviors do not operate on logic, and that the attempt to force them into a logical mold is not understanding but misrecognition — a sustained and socially sanctioned confusion of the shadow for the substance, of the narrative for the loop, of the mask for the face.

CHAPTER NINE

Conclusion: The Seam We Cannot See

In the end, the most difficult truth is not that we are less rational than we suppose. It is that our irrationality is not a failure of reason but a success of camouflage — a structural adaptation that allows affect-driven creatures to inhabit a world that demands the performance of rationality. We do not reason and then act. We loop, and then we explain. And the explanation is so compelling, so seamlessly woven into the fabric of our experience, that we cannot see the seam. The loop is invisible because it is the ground; the reason is visible because it is the figure we paint upon it.

The argument of this monograph has been architectural rather than debunking. It has not denied that human beings reason, deliberate, plan, and sometimes act from reasons genuinely held in advance. It has denied the reconstructive paradigm's central inference — from the availability of a reason to the causal priority of that reason — and has shown that across five independent research programs, the inference fails in exactly the systematic way that an interpreter-plus-loop architecture predicts. The narrator is real, and tireless, and sincere; and his sincerity is the very thing that must not be taken as evidence.

What remains, once the inversion is accepted, is not cynicism but a different and more demanding form of charity. To understand another person on the dynamical picture is to give up the flattering fiction that one could simply swap in better arguments — and to take up instead the harder work of attending to basins and couplings, to the environments and rhythms and histories that shaped them. It is to recognize that the defendant, the partisan, the beloved, and the self are all of the same metaphysical kind: dynamical systems wearing the mask of reasoners, because the mask is what the social world requires and because, after long wearing, the mask is what the mirror shows.

To understand this is to understand something not merely about the limits of human rationality, but about the invention of rationality itself — as a post-hoc artifact of a creature who must live as if he were the author of his own conduct, even as he is authored by the loops he cannot see.

References

1. Davidson, D. (1963). Actions, reasons, and causes. *Journal of Philosophy*, 60(23), 685–700.
2. Lewis, M. D. (2000). Emotional self-organization at three time scales. In M. D. Lewis & I. Granic (Eds.), *Emotion, Development, and Self-Organization: Dynamic Systems Approaches to Emotional Development* (pp. 37–69). Cambridge University Press.
3. Lewis, M. D., & Granic, I. (Eds.). (2000). *Emotion, Development, and Self-Organization: Dynamic Systems Approaches to Emotional Development*. Cambridge University Press.
4. Nisbett, R. E., & Wilson, T. D. (1977). Telling more than we can know: Verbal reports on mental processes. *Psychological Review*, 84(3), 231–259.
5. Gazzaniga, M. S. (1995). Principles of human brain organization derived from split-brain studies. *Neuron*, 14(2), 217–228.
6. Gazzaniga, M. S. (2000). Cerebral specialization and interhemispheric communication: Does the corpus callosum enable the human condition? *Brain*, 123(7), 1293–1326.
7. Gazzaniga, M. S. (2024). Michael Gazzaniga (PhD '65), Neuroscientist and Pioneer in Split Brain Research [Interview]. *Caltech Heritage Project*. <https://heritageproject.caltech.edu/interviews-updates/michael-gazzaniga>
8. Johansson, P., Hall, L., Sikström, S., & Olsson, A. (2005). Failure to detect mismatches between intention and outcome in a simple decision task. *Science*, 310(5745), 116–119.
9. Johansson, P., Hall, L., Sikström, S., Tärning, B., & Lind, A. (2006). How something can be said about telling more than we can know: On choice blindness and introspection. *Consciousness and Cognition*, 15(4), 673–692.
10. Wegner, D. M., & Wheatley, T. (1999). Apparent mental causation: Sources of the experience of will. *American Psychologist*, 54(7), 480–492.
11. Wegner, D. M. (2002). *The Illusion of Conscious Will*. MIT Press.
12. Haidt, J. (2001). The emotional dog and its rational tail: A social intuitionist approach to moral judgment. *Psychological Review*, 108(4), 814–834.
13. Hirstein, W. (2005). *Brain Fiction: Self-Deception and the Riddle of Confabulation*. MIT Press.
14. Mercier, H., & Sperber, D. (2011). Why do humans reason? Arguments for an argumentative theory. *Behavioral and Brain Sciences*, 34(2), 57–74.
15. Nahmias, E. (2002). When consciousness matters: A critical review of Daniel Wegner's *The Illusion of Conscious Will*. *Philosophical Psychology*, 15(4), 527–541.
16. Damasio, A. R. (1994). *Descartes' Error: Emotion, Reason, and the Human Brain*. Putnam.
17. Paxton, J. M., & Greene, J. D. (2010). Moral reasoning: Hints and allegations. *Topics in Cognitive Science*, 2(3), 511–527.
18. Bloom, P. (2010). How do morals change? *Nature*, 464(7288), 490.
19. Haidt, J., & Bjorklund, F. (2008). Social intuitionists answer six questions about moral psychology. In W. Sinnott-Armstrong (Ed.), *Moral Psychology, Vol. 2: The Cognitive Science of Morality* (pp. 181–217). MIT Press.

The Retroactive Architecture of Agency — Expanded Monograph Edition — 2026